

```

1 #!usr/bin/python
2 #-*-encoding:UTF-8-*-
3 #Date:2021/10/25
4 #Author:Dasein
5 #载入包：设置;载入数据
6 import pandas as pd
7 import numpy as np
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 sns.set(style='darkgrid',font_scale=1.3)
11 plt.rcParams['font.family']='SimHei'
12 plt.rcParams['axes.unicode_minus']=False
13 import warnings
14 warnings.filterwarnings('ignore')

```

```

1 df1 = pd.read_csv("E:\\dasein_py\\Data Analysis\\Telecommunication_da\\WA_Fn-UseC_-Telco-Customer-Churn.csv")

```

```

1 print(df1.info())
2 print(df1.shape)

```

```

1 <class 'pandas.core.frame.DataFrame'>
2 RangeIndex: 7043 entries, 0 to 7042
3 Data columns (total 21 columns):
4  #   Column                Non-Null Count  Dtype
5  ---  ---
6  0   customerID            7043 non-null   object
7  1   gender                 7043 non-null   object
8  2   SeniorCitizen          7043 non-null   int64
9  3   Partner                7043 non-null   object
10  4   Dependents             7043 non-null   object
11  5   tenure                 7043 non-null   int64
12  6   PhoneService           7043 non-null   object
13  7   MultipleLines          7043 non-null   object
14  8   InternetService        7043 non-null   object
15  9   OnlineSecurity         7043 non-null   object
16  10  OnlineBackup            7043 non-null   object
17  11  DeviceProtection       7043 non-null   object
18  12  TechSupport            7043 non-null   object
19  13  StreamingTV            7043 non-null   object
20  14  StreamingMovies        7043 non-null   object
21  15  Contract               7043 non-null   object
22  16  PaperlessBilling       7043 non-null   object
23  17  PaymentMethod          7043 non-null   object
24  18  MonthlyCharges         7043 non-null   float64
25  19  TotalCharges           7043 non-null   object
26  20  Churn                  7043 non-null   object
27 dtypes: float64(1), int64(2), object(18)
28 memory usage: 1.1+ MB
29 None
30 (7043, 21)

```

```

1 quantative = [i for i in df1.columns if df1[i].dtype!=object]
2 quanlitive = [i for i in df1.columns if df1[i].dtype==object]
3 print("Quantative counts:{}, Quanlitive counts:{}".format(len(quantative),len(quanlitive)))

```

```

1 Quantative counts:3, Quanlitive counts:18

```

Data Overall

- Dtype: float64 & string (Quantative:3, Quanlitive:18)
- Case counts: 7043
- Variable counts: 21

```

1 df1.describe()

```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

```
1 df1.columns
```

```
1 Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
2       'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
3       'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
4       'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
5       'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
6       dtype='object')
```

Variable Notes

- customerID: ID
- gender
- SeniorCitizen: Whether the customer is a senior citizen or not (1, 0)
- Partner: Whether the customer has a partner or not (Yes, No)
- Dependents: Whether the customer has dependents or not (Yes, No)
- tenure: Number of months the customer has stayed with the company
- PhoneService: Whether the customer has a phone service or not (Yes, No)
- MultipleLines: Whether the customer has multiple lines or not (Yes, No, No phone service)
- InternetService: Customer's internet service provider (DSL, Fiber optic, No)
- OnlineSecurity: Whether the customer has online security or not (Yes, No, No internet service)
- OnlineBackup: Whether the customer has online backup or not (Yes, No, No internet service)
- DeviceProtection: Whether the customer has device protection or not (Yes, No, No internet service)
- TechSupport: Whether the customer has tech support or not (Yes, No, No internet service)
- StreamingTV: Whether the customer has streaming TV or not (Yes, No, No internet service)
- StreamingMovies: Whether the customer has streaming movies or not (Yes, No, No internet service)
- Contract: The contract term of the customer (Month-to-month, One year, Two year)
- PaperlessBilling: Whether the customer has paperless billing or not (Yes, No)
- PaymentMethod: The customer's payment method (Electronic check, Mailed check, Bank transfer (automatic), Credit card (automatic))
- MonthlyCharges: The amount charged to the customer monthly
- TotalCharges: The total amount charged to the customer
- Churn: Whether the customer churned or not (Yes or No)

```
1 df1.head(3)
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...

3 rows × 21 columns

Insights

- customerID可以drop
- Quantative中SeniorCitizen是0-1变量
- Quanlitive数据需要重编码

```
1 df1.drop('customerID',axis=1,inplace=True) #drop colName: CustomerID
```

```
1 df1.head(3) #double check data after drop
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	De
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	No
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	Yes
2	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	No

```
1 total = df1.isnull().sum()
2 null_percentage = total/df1.isnull().count()
3 null_percentage
```

1	gender	0.0
2	SeniorCitizen	0.0
3	Partner	0.0
4	Dependents	0.0
5	tenure	0.0
6	PhoneService	0.0
7	MultipleLines	0.0
8	InternetService	0.0
9	OnlineSecurity	0.0
10	OnlineBackup	0.0
11	DeviceProtection	0.0
12	TechSupport	0.0
13	StreamingTV	0.0
14	StreamingMovies	0.0

```
15 Contract          0.0
16 PaperlessBilling  0.0
17 PaymentMethod     0.0
18 MonthlyCharges    0.0
19 TotalCharges      0.0
20 Churn             0.0
21 dtype: float64
```

Insights

- Hypothesis: probable duplicates.

```
1 print(df1.duplicated().sum())
2 df1=df1.drop_duplicates(subset=None, keep='first',inplace=False)
```

```
1 22
```

```
1 #double check去重之后data
2 print(df1.duplicated().sum())
```

```
1 0
```

```
1 # TotalCharges应该是数值型，需要强制类型转换
2 # df1['TotalCharges']=df1['TotalCharges'].astype('float64')
3 df1['TotalCharges'] = df1['TotalCharges'].apply(pd.to_numeric, errors='coerce')
```

```
1 df1['TotalCharges'].dtype
```


```
1 dtype('float64')
```

```
1 df1['TotalCharges'].isnull().sum() #有缺失值
```

```
1 11
```

```
1 print('TotalCharges数据分布')
2 plt.figure(figsize=(14,5))
3 plt.plot(color='#003380')
4 #1
5 plt.subplot(1,3,1)
6 plt.title("TotalCharges distplot")
7 sns.distplot(df1.TotalCharges)
8 #2
9 plt.subplot(1,3,2)
10 plt.title("Churn Distplot")
11 sns.distplot(df1[df1.Churn=='No']['TotalCharges'])
12 #2
13 plt.subplot(1,3,3)
14 plt.title("Churn + TotalCharges Distplot")
15 sns.distplot(df1[df1['Churn']=='Yes']['TotalCharges'])
16 plt.show()
```

```
1 TotalCharges数据分布
```



Insights

- TotalCharges* 偏态分布，需要用中值填充缺失值。

```
1 df1.fillna({'TotalCharges':df1.TotalCharges.median()},inplace=True)
```

```
1 df1.TotalCharges.isnull().sum() #已经没有缺失值
```

```
1 0
```

```
1 #重编码 'Churn'，定性转定量的哑变量
2 #df1.Churn=df1.Churn.map({'Yes':1,'No':0})
3 df1.Churn.replace(to_replace='Yes',value=1,inplace=True)
4 df1.Churn.replace(to_replace='No',value=0,inplace=True)
5 df1.Churn.isnull().sum()
```

```
1 0
```


```
1 df1.Churn.describe()
```

```
1 count    7021.000000
2 mean      0.264492
3 std       0.441094
4 min       0.000000
5 25%       0.000000
6 50%       0.000000
7 75%       1.000000
8 max       1.000000
9 Name: Churn, dtype: float64
```


Insights

- 平均流失率 26.45%。

```
1 Churn_Count=df1.Churn.value_counts()
2 Churn_Lab=df1.Churn.value_counts().index
3 plt.figure(figsize=(5,5))
4 plt.pie(Churn_Count,labels=Churn_Lab,
5         colors=["#00338D","red"],
6         explode=(0.3,0),
7         autopct="%1.1f%%",
8         shadow=True)
9 plt.title("Customer Churn Pie Chart")
10 plt.show()
```



```
1 plt.figure(figsize=(4,6))
2 plt.plot(color='#00338D')
3 fig = sns.boxplot(x="Churn",y="tenure",data=df1)
4 plt.title("tenure - Churn Boxplot")
5 fig.axis(ymin=0,ymax=80)
6 plt.show()
```



- tenure越小流失率越显著


```
1 df2 = df1.apply(lambda x:pd.factorize(x)[0]) #转换成因子
2 df2.head(5)
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	De
0	0	0	0	0	0	0	0	0	0	0	0
1	1	0	1	0	1	1	1	0	1	1	1
2	1	0	1	0	2	1	1	0	1	0	0
3	1	0	1	0	3	0	0	0	1	1	1
4	0	0	1	0	2	1	1	1	0	1	0

```
1 var = list(df2.columns)
2 var.remove("Churn")
3 var.remove("tenure")
4 var.remove("MonthlyCharges")
5 var.remove("TotalCharges")
6 plt.figure(figsize=(30,25))
```

```
7 a=0
8 for item in var:
9     a+=1
10    plt.subplot(4,5,a)
11    plt.title('Barplot by '+ item)
12    sns.countplot(x=item,data=df2,
13                  color="#00338D")
14 #sns.countplot(x=None, y=None,
15 #hue=None, data=None, order=None,
16 #hue_order=None, orient=None, color=None,
17 #palette=None, saturation=0.75, dodge=True, ax=None, **kwargs)
```



Insights

- gender对Churn的影响不显著

```
1 df2.drop("gender",axis=1,inplace=True)
```

```
1 -----
2
3 KeyError                                Traceback (most recent call last)
4
5 <ipython-input-28-67322b8776aa> in <module>
6 ----> 1 df2.drop("gender",axis=1,inplace=True)
```

```
1 D:\anaconda\lib\site-packages\pandas\core\frame.py in drop(self, labels, axis, index, columns, level, inplace, errors)
2     3988         weight 1.0    0.8
3     3989         """
4     -> 3990         return super().drop(
5     3991             labels=labels,
6     3992             axis=axis,
```

```
1 D:\anaconda\lib\site-packages\pandas\core\generic.py in drop(self, labels, axis, index, columns, level, inplace, errors)
2     3934         for axis, labels in axes.items():
3     3935             if labels is not None:
4     -> 3936                 obj = obj._drop_axis(labels, axis, level=level, errors=errors)
5     3937
6     3938         if inplace:
```

```
1 D:\anaconda\lib\site-packages\pandas\core\generic.py in _drop_axis(self, labels, axis, level, errors)
2     3968         new_axis = axis.drop(labels, level=level, errors=errors)
3     3969     else:
4 -> 3970         new_axis = axis.drop(labels, errors=errors)
5     3971     result = self.reindex(**{axis_name: new_axis})
6     3972
```

```
1 D:\anaconda\lib\site-packages\pandas\core\indexes\base.py in drop(self, labels, errors)
2     5016         if mask.any():
3     5017             if errors != "ignore":
4 -> 5018                 raise KeyError(f"{labels[mask]} not found in axis")
5     5019             indexer = indexer[~mask]
6     5020         return self.delete(indexer)
```

```
1 KeyError: "['gender'] not found in axis"
```

```
1 df2.isnull().sum() #转换成因子之后没有缺失值，不需要fillna（TotalCharges已经填充缺失值）
```


```
1 SeniorCitizen      0
2 Partner            0
3 Dependents         0
4 tenure            0
5 PhoneService       0
6 MultipleLines      0
7 InternetService    0
8 OnlineSecurity     0
9 OnlineBackup       0
10 DeviceProtection  0
11 TechSupport       0
12 StreamingTV       0
13 StreamingMovies   0
14 Contract          0
15 PaperlessBilling  0
16 PaymentMethod     0
17 MonthlyCharges    0
18 TotalCharges      0
19 Churn             0
20 dtype: int64
```

```
1 corr = df2.corr()
2 corr
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```


	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBack
SeniorCitizen	1.000000	-0.015553	-0.212115	0.009452	0.008909	0.113409	-0.031221	-0.210370	-0.143900
Partner	-0.015553	1.000000	-0.451659	-0.098113	-0.018728	-0.115992	-0.003667	-0.084330	0.087952
Dependents	-0.212115	-0.451659	1.000000	0.045761	-0.001092	-0.020715	0.046608	0.192658	0.065093
tenure	0.009452	-0.098113	0.045761	1.000000	-0.017391	0.061467	-0.007970	0.020317	-0.060309
PhoneService	0.008909	-0.018728	-0.001092	-0.017391	1.000000	0.675973	0.387549	0.125042	0.129032
MultipleLines	0.113409	-0.115992	-0.020715	0.061467	0.675973	1.000000	0.188826	-0.065972	-0.129333
InternetService	-0.031221	-0.003667	0.046608	-0.007970	0.387549	0.188826	1.000000	0.606107	0.649514
OnlineSecurity	-0.210370	-0.084330	0.192658	0.020317	0.125042	-0.065972	0.606107	1.000000	0.620365
OnlineBackup	-0.143900	0.087952	0.065093	-0.060309	0.129032	-0.129333	0.649514	0.620365	1.000000
DeviceProtection	-0.156410	-0.096813	0.158328	0.040275	0.138544	-0.012102	0.661669	0.747520	0.600141
TechSupport	-0.223293	-0.071483	0.182923	0.036426	0.123035	-0.065817	0.608130	0.789952	0.615611
StreamingTV	-0.129375	-0.082304	0.142145	0.029948	0.171477	0.031247	0.711946	0.700176	0.602861
StreamingMovies	-0.120015	-0.077925	0.127508	0.034361	0.165127	0.029227	0.708061	0.703203	0.605631
Contract	-0.143624	-0.293042	0.241912	0.119246	0.003101	0.082152	0.102456	0.393394	0.038225
PaperlessBilling	-0.156196	-0.015776	0.112220	-0.009923	-0.016824	-0.133094	0.137056	0.332537	0.259546
PaymentMethod	-0.094091	-0.131842	0.122957	0.073367	-0.003547	0.024891	0.008899	0.214518	0.004219
MonthlyCharges	0.048736	-0.034681	-0.030433	0.039656	-0.141515	0.023609	-0.289498	-0.220075	-0.283567
TotalCharges	0.022996	-0.040026	0.004450	0.104648	-0.028946	0.013971	-0.035305	-0.023596	-0.051101
Churn	0.151619	0.149135	-0.163459	-0.142337	0.011323	0.037429	-0.047366	-0.333144	-0.075052

```
1 plt.figure(figsize=(15,15))
2 sns.set(font_scale=1.25)
3 ax=sns.heatmap(corr,
4                 xticklabels=corr.columns,
5                 yticklabels=corr.columns,
6                 linewidths=0.6,annot=True,
7                 cbar=True,cmap="rainbow",fmt='.4f',
8                 annot_kws={'size': 10})
9 plt.title("Corr Heatmap")
10 plt.savefig("Corr Heatmap.png",dpi=100)
11 plt.show()
```



Insights

- 极强相关变量
- **MultipleLines - PhoneService**之间有很强共线性。
- 相关系数矩阵中心的变量之间具有极强的相关性（共线性）
 - *OnlineSecurity / InternetService / OnlineBackup / DeviceProtection / TechSupport / StreamingTV / StreamingMovies*
- 没有与Churn具有极强相关性的变量。
 - *TotalCharges / MonthlyCharges / OnlineBackup / InternetService / MultipleLines / PhoneService* 与Churn相关系数极小。
 - *TotalCharges*与其他变量相关系数均很小。

热力图效果不是很显著。

```
1 #独热编码
2 df_onehot = pd.get_dummies(df1.iloc[:, :])
3 df_onehot
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_No	Partner_Yes	Dependents_No
0	0	1	29.85	29.85	0	1	0	0	1	1
1	0	34	56.95	1889.50	0	0	1	1	0	1
2	0	2	53.85	108.15	1	0	1	1	0	1
3	0	45	42.30	1840.75	0	0	1	1	0	1
4	0	2	70.70	151.65	1	1	0	1	0	1
...
7038	0	24	84.80	1990.50	0	0	1	0	1	0
7039	0	72	103.20	7362.90	0	1	0	0	1	0
7040	0	11	29.60	346.45	0	1	0	0	1	0
7041	1	4	74.40	306.60	1	0	1	0	1	1
7042	0	66	105.65	6844.50	0	0	1	1	0	1

7021 rows × 46 columns

- 独热编码的相关系数细分将变量数值拆分成子变量，研究自变量和Churn之间的相关性大小，进一步研究变量与Churn的相关性。
- gender和phoneservice不相关，所以继续drop phoneservice，用drop之后的var进行卡方检验频数比较。

```
1 df1.drop("PhoneService",axis=1,inplace=True)
```

```
1 var.remove('PhoneService')
```

```
1 var.remove('gender')
```

```
1 -----
2
3 ValueError                                Traceback (most recent call last)
4
5 <ipython-input-35-3cb360e98008> in <module>
6 ----> 1 var.remove('gender')
```

```
1 ValueError: list.remove(x): x not in list
```

```
1 #交叉分析
2 print('kf_var与Churn的进行交叉分析','\n')
3 for item in var:
4     print("-----Churn by {}-----".format(item))
5     print(pd.crosstab(df2.Churn,df2[item],normalize=0),'\n')
```

```
1 kf_var与Churn的进行交叉分析
2
3 -----Churn by SeniorCitizen-----
4 SeniorCitizen      0      1
5 Churn
6 0      0.871030  0.128970
7 1      0.744211  0.255789
8
9 -----Churn by Partner-----
10 Partner      0      1
11 Churn
12 0      0.529241  0.470759
13 1      0.360258  0.639742
14
15 -----Churn by Dependents-----
16 Dependents      0      1
17 Churn
18 0      0.654531  0.345469
19 1      0.824448  0.175552
20
21 -----Churn by MultipleLines-----
22 MultipleLines      0      1      2
23 Churn
24 0      0.099148  0.490124  0.410728
25 1      0.091546  0.450727  0.457728
```

```

26
27 -----Churn by InternetService-----
28 InternetService      0      1      2
29 Churn
30 0                    0.379938  0.348373  0.271689
31 1                    0.246096  0.695207  0.058697
32
33 -----Churn by OnlineSecurity-----
34 OnlineSecurity      0      1      2
35 Churn
36 0                    0.394462  0.333850  0.271689
37 1                    0.782445  0.158858  0.058697
38
39 -----Churn by OnlineBackup-----
40 OnlineBackup      0      1      2
41 Churn
42 0                    0.369094  0.359218  0.271689
43 1                    0.281637  0.659666  0.058697
44
45 -----Churn by DeviceProtection-----
46 DeviceProtection    0      1      2
47 Churn
48 0                    0.364833  0.363478  0.271689
49 1                    0.647819  0.293484  0.058697
50
51 -----Churn by TechSupport-----
52 TechSupport      0      1      2
53 Churn
54 0                    0.392525  0.335786  0.271689
55 1                    0.774367  0.166936  0.058697
56
57 -----Churn by StreamingTV-----
58 StreamingTV      0      1      2
59 Churn
60 0                    0.361735  0.366576  0.271689
61 1                    0.502962  0.438341  0.058697
62
63 -----Churn by StreamingMovies-----
64 StreamingMovies    0      1      2
65 Churn
66 0                    0.357668  0.370643  0.271689
67 1                    0.500808  0.440495  0.058697
68
69 -----Churn by Contract-----
70 Contract      0      1      2
71 Churn
72 0                    0.427963  0.253098  0.318939
73 1                    0.884760  0.089391  0.025848
74
75 -----Churn by PaperlessBilling-----
76 PaperlessBilling    0      1
77 Churn
78 0                    0.536406  0.463594
79 1                    0.749058  0.250942
80
81 -----Churn by PaymentMethod-----
82 PaymentMethod      0      1      2      3
83 Churn
84 0                    0.250581  0.250581  0.249032  0.249806
85 1                    0.573506  0.162628  0.138934  0.124933

```

- Crosstab中若变量取值对应的Churn - Yes的百分比差异越大，说明该变量对Churn - Yes的影响越显著。
 - SeniorCitizen：在年轻用户流失、留存的占比都很高。
 - Partner：单身越流失。
 - Dependents：经济不独立越流失。
 - StreamingMovies/StreamTvs/Multiplelines：不显著。
 - InternetService：Fiber Optic更易流失。
 - OnlineSecurity/OnlineBackup/DeviceProtection/TechSupport：没开通容易流失。
 - Contract：逐月订阅易流失。
 - Check：电子支票易流失。

```

1 from scipy import stats
2 def ANOVA(x):
3     index_list = list(df2['Churn'].value_counts().keys())
4     args=[]
5     for i in index_list:

```

```
6         args.append(df2[df2['Churn']==i][x])
7     w,p=stats.levene(*args) #齐性检验
8     if p < 0.05:
9         print('Churn By {}, p值是{:.2f}, 小于0.05, 表明方差齐性检验不通过, 不可做方差分析.'.format(x,p),'\n')
10    else:
11        f,p_value = stats.f_oneway(*args)#方差检验
12        print('Churn By {},f值是{:.2f}, p值是{:.2f}'.format(x,f,p_value),'\n')
13        if p_value <0.05:
14            print("Churn by {}有显著性差异, 可进行均值比较.".format(x),'\n')
15        else:
16            print("Churn by {}没有显著性差异, 不可进行均值比较.".format(x),'\n')
17
```

```
1 print("MonthlyCharges和TotalCharges齐性检验和方差分析, 如下: ",'\n')
2 ANOVA('MonthlyCharges')
3 ANOVA('TotalCharges')
```

```
1 MonthlyCharges和TotalCharges齐性检验和方差分析, 如下:
2
3 Churn By MonthlyCharges,f值是3.34, p值是0.07。
4
5 Churn by MonthlyCharges没有显著性差异, 不可进行均值比较。
6
7 Churn By TotalCharges,f值是5.13, p值是0.02。
8
9 Churn by TotalCharges有显著性差异, 可进行均值比较。
```

```
1 df1[["MonthlyCharges","TotalCharges"]]
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	MonthlyCharges	TotalCharges
0	29.85	29.85
1	56.95	1889.50
2	53.85	108.15
3	42.30	1840.75
4	70.70	151.65
...
7038	84.80	1990.50
7039	103.20	7362.90
7040	29.60	346.45
7041	74.40	306.60
7042	105.65	6844.50

7021 rows × 2 columns

- MonthlyCharges & TotalCharges 量纲差异大。
- gender, id, PhoneService对Churn影响不显著, 应该drop。

```
1 #标准化
2 from sklearn.preprocessing import StandardScaler
3 #sklearn.preprocessing.StandardScaler(copy=True, with_mean=True, with_std=True)
4 scaler = StandardScaler(copy=False)
5 scaler.fit_transform(df1[['MonthlyCharges','TotalCharges']]) #拟合数据
6 df1[['MonthlyCharges','TotalCharges']]=scaler.transform(df1[['MonthlyCharges','TotalCharges']]) #数据标准化
7 df1[['MonthlyCharges','TotalCharges']].head()
```



```

1 import sklearn #特征工程
2 from sklearn import preprocessing #数据预处理
3 from sklearn.preprocessing import LabelEncoder #编码转换
4 def labelencoder(x):
5     df1[x]=LabelEncoder().fit_transform(df1[x])
6 for i in range(len(df_obj.columns)):
7     labelencoder(df_obj.columns[i])
8 print(list(map(Labs,df_obj.columns)))

```

```

1 gender -- [0 1]
2 Partner -- [1 0]
3 Dependents -- [0 1]
4 MultipleLines -- [0 1]
5 InternetService -- [0 1 2]
6 OnlineSecurity -- [0 1]
7 OnlineBackup -- [1 0]
8 DeviceProtection -- [0 1]
9 TechSupport -- [0 1]
10 StreamingTV -- [0 1]
11 StreamingMovies -- [0 1]
12 Contract -- [0 1 2]
13 PaperlessBilling -- [1 0]
14 PaymentMethod -- [2 3 0 1]
15 [None, None, None, None, None, None, None, None, None, None, None, None, None, None, None]

```

```

1 list(map(Labs,df1.columns))

```

```

1 gender -- [0 1]
2 SeniorCitizen -- [0 1]
3 Partner -- [1 0]
4 Dependents -- [0 1]
5 tenure -- [ 1 34  2 45  8 22 10 28 62 13 16 58 49 25 69 52 71 21 12 30 47 72 17 27
6          5 46 11 70 63 43 15 60 18 66  9  3 31 50 64 56  7 42 35 48 29 65 38 68
7          32 55 37 36 41  6  4 33 67 23 57 61 14 20 53 40 59 24 44 19 54 51 26  0
8          39]
9 MultipleLines -- [0 1]
10 InternetService -- [0 1 2]
11 OnlineSecurity -- [0 1]
12 OnlineBackup -- [1 0]
13 DeviceProtection -- [0 1]
14 TechSupport -- [0 1]
15 StreamingTV -- [0 1]
16 StreamingMovies -- [0 1]
17 Contract -- [0 1 2]
18 PaperlessBilling -- [1 0]
19 PaymentMethod -- [2 3 0 1]
20 MonthlyCharges -- [-1.16413536 -0.26281076 -0.36591432 ... -0.05826662 -0.68686569
21                   0.46057706]
22 TotalCharges -- [-0.99733366 -0.17635202 -0.96276648 ... -0.85756393 -0.87515655
23                  2.01113704]
24 Churn -- [0 1]

```

```

1 [None,
2   None,
3   None,
4   None,
5   None,
6   None,
7   None,
8   None,
9   None,
10  None,
11  None,
12  None,
13  None,
14  None,
15  None,
16  None,
17  None,
18  None,
19  None]

```

```
1 # #处理样本不平衡，分拆变量
2 # df1.drop("gender",axis=1,inplace=True)
3 # df1.drop("PhoneService",axis=1,inplace=True)
```

```
1 x=df1[var]
2 y=df1['Churn'].values
3 x
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	SeniorCitizen	Partner	Dependents	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	...
0	0	1	0	0	0	0	1	0	0	0
1	0	0	0	0	0	1	0	1	0	0
2	0	0	0	0	0	1	1	0	0	0
3	0	0	0	0	0	1	0	1	1	0
4	0	0	0	0	1	0	0	0	0	0
...
7038	0	1	1	1	0	1	0	1	1	1
7039	0	1	1	1	1	0	1	1	0	1
7040	0	1	1	0	0	1	0	0	0	0
7041	1	1	0	1	1	0	0	0	0	0
7042	0	0	0	0	1	1	0	1	1	0

7021 rows × 14 columns

```
1 from sklearn.model_selection import StratifiedShuffleSplit #分层抽样
2 from sklearn.model_selection import train_test_split #数据集训练集划分
```

```
1 #分层抽样stratified random sampling、过抽样、欠抽样，抽样上面多试错
2 sss=StratifiedShuffleSplit(n_splits=5,test_size=.2,random_state=0)
3 print(sss)
4 print(sss.split(x,y))
```

```
1 StratifiedShuffleSplit(n_splits=5, random_state=0, test_size=0.2,
2     train_size=None)
3 <generator object BaseShuffleSplit.split at 0x000001EF6EA42900>
```

```
1 print("训练数据和测试数据被分成的份数：",sss.get_n_splits(x,y))
2 #拆分训练集和测试集
3 for train_index,test_index,in sss.split(x,y):
4     print("train:",train_index,"test:",test_index)
5     x_train,x_test=x.iloc[train_index],x.iloc[test_index]
6     y_train,y_test= y[train_index],y[test_index]
```

```
1 训练数据和测试数据被分成的份数： 5
2 train: [5297 5907 3429 ... 4096 6084 3612] test: [4979 2569 5247 ... 1572 4876 4997]
3 train: [4203 5971 767 ... 1505 230 4637] test: [3201 692 688 ... 4736 3769 5207]
4 train: [5070 2818 1921 ... 4575 6509 1607] test: [1213 3852 1396 ... 1855 2852 1846]
5 train: [1468 2332 1900 ... 6038 5207 943] test: [5733 3682 4429 ... 6390 944 6816]
6 train: [5861 1463 4124 ... 4026 6659 4286] test: [6811 5731 3968 ... 2 4805 6708]
```

```
1 print("分层抽样数据特征：",x.shape,"train特征:",x_train_.shape,"test特征：",x_test_.shape)
2 print("分层抽样数据特征：",y.shape,"train特征:",y_train_.shape,"test特征：",y_test_.shape)
```



```
1  分层抽样数据特征： (7021, 14) train特征： (5616, 14) test特征： (1405, 14)
2  分层抽样数据特征： (7021,) train特征： (5616,) test特征： (1405,)
```

```
1  # sklearn.linear_model.RidgeCV(alphas=(0.1, 1.0, 10.0), fit_intercept=True, normalize=False, scoring=None, cv=None, gcv_mode=None,
store_cv_values=False)
2  from sklearn.linear_model import RidgeClassifier, RidgeCV # 岭回归
3  from sklearn.metrics import accuracy_score
4  rcv = RidgeClassifier()
5  rcv.fit(x_train,y_train)
6  pred = rcv.predict(x_test)
7  print(accuracy_score(y_test,pred))
```

```
1  from sklearn.model_selection import train_test_split #数据集训练集划分
2  x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.2,random_state=22)
```

```
1  # 训练分类模型
2  from sklearn import metrics
3  from sklearn.metrics import recall_score
4  from sklearn.metrics import accuracy_score
5  from sklearn.metrics import precision_score
6  from sklearn.metrics import f1_score
7  from sklearn.ensemble import RandomForestClassifier #随机森林
8  from sklearn.svm import SVC#支持向量机
9  from sklearn.linear_model import LogisticRegression #逻辑回归
10 from sklearn.neighbors import KNeighborsClassifier #k邻近算法
11 from sklearn.naive_bayes import GaussianNB #朴素贝叶斯
12 from sklearn.tree import DecisionTreeClassifier #决策树
13 from sklearn.ensemble import AdaBoostClassifier #分类器算法
14 from sklearn.ensemble import GradientBoostingClassifier #梯度提升
15 from xgboost import XGBClassifier
16 from catboost import CatBoostClassifier
17 from sklearn.linear_model import RidgeClassifier # 岭
18 from sklearn.neural_network import MLPClassifier #神经网络
19 from sklearn.linear_model import SGDClassifier
20 from sklearn.ensemble import BaggingClassifier
21 from sklearn.ensemble import ExtraTreesClassifier
22 from xgboost import XGBClassifier
23 import time
```

```
1  Classifiers = [{"Random Forest",RandomForestClassifier()},
2                  ["Support Vector Machine",SVC()],
3                  ["LogisticRegression",LogisticRegression()],
4                  ["KNeighbor",KNeighborsClassifier(n_neighbors=5)],
5                  ["Naive Bayes",GaussianNB()],
6                  ["Decision Tree",DecisionTreeClassifier()],
7                  ["GradientBoostingClassifier",GradientBoostingClassifier()],
8                  ["XGB",XGBClassifier()],
9                  ["CatBoost",CatBoostClassifier(logging_level='silent')],
10                 ['RidgeClassifier',RidgeClassifier()],
11                 ['MLPClassifier',MLPClassifier(solver='lbfgs',activation = 'tanh',
12                 max_iter = 50,alpha = 0.001,
13                 hidden_layer_sizes = (10,30),
14                 random_state = 1,verbose = True)],
15                 ['SGDClassifier',SGDClassifier()],
16                 ['XGBClassifier',XGBClassifier()],
17                 ['BaggingClassifier',BaggingClassifier()],
18                 ['XGBClassifier',XGBClassifier()]
19  ]
```

```
1  import time
2  Classify_result=[]
3  names=[]
4  prediction=[]
5  for name,classifier in Classifiers:
6      classifier=classifier
7      t1 = time.time()
8      classifier.fit(x_train,y_train)
9      y_pred=classifier.predict(x_test)
10     t2=time.time()
11     precision=precision_score(y_test,y_pred)
12     f1score = f1_score(y_test, y_pred)
13     time_diff = t2 -t1
14     class_eva=pd.DataFrame([precision,f1score,time_diff])
15     Classify_result.append(class_eva)
16     name=pd.Series(name)
17     names.append(name)
```

```
18 | y_pred=pd.Series(y_pred)
19 | prediction.append(y_pred)
```

```
1 | [16:03:33] WARNING: D:\Build\xgboost\xgboost-1.4.2.git\src\learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with
the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
2 | [16:03:39] WARNING: D:\Build\xgboost\xgboost-1.4.2.git\src\learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with
the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
3 | [16:03:39] WARNING: D:\Build\xgboost\xgboost-1.4.2.git\src\learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with
the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
```

```
1 | names = pd.DataFrame(names)
2 | names=names[0].tolist()
3 | names
```

```
1 | ['Random Forest',
2 |  'Support Vector Machine',
3 |  'LogisticRegression',
4 |  'KNeighbor',
5 |  'Naive Bayes',
6 |  'Decision Tree',
7 |  'GradientBoostingClassifier',
8 |  'XGB',
9 |  'CatBoost',
10 | 'RidgeClassifier',
11 | 'MLPClassifier',
12 | 'SGDClassifier',
13 | 'XGBClassifier',
14 | 'BaggingClassifier',
15 | 'XGBClassifier']
```

```
1 | result = pd.concat(Classify_result,axis=1)
```

```
1 | result.columns =names
```

```
1 | result.index = ["precision",'f1score',"time_diff"]
```

```
1 | result.T
```

```
1 | .dataframe tbody tr th {
2 |     vertical-align: top;
3 | }
4 |
5 | .dataframe thead th {
6 |     text-align: right;
7 | }
```

	precision	f1score	time_diff
Random Forest	0.546032	0.500728	0.553662
Support Vector Machine	0.624000	0.501608	1.780465
LogisticRegression	0.578544	0.477093	0.029593
KNeighbor	0.506702	0.507383	0.178499
Naive Bayes	0.507937	0.584475	0.005983
Decision Tree	0.476584	0.470748	0.012998
GradientBoostingClassifier	0.596215	0.548621	0.435386
XGB	0.550152	0.516405	0.482712
CatBoost	0.579618	0.530612	4.346667
RidgeClassifier	0.614973	0.411449	0.010227
MLPClassifier	0.612903	0.557185	0.662815
SGDClassifier	0.736842	0.136585	0.026943
XGBClassifier	0.550152	0.516405	0.475735
BaggingClassifier	0.526946	0.498584	0.144153
XGBClassifier	0.550152	0.516405	0.473733